

Big Data Project

LoudAcre Mobile Kmeans Clustering

```
import org.apache.spark.mllib.linalg.Vectors
import org.apache.spark.mllib.clustering.KMeans
import org.apache.spark.sql.functions._
```

- Loading the Dataset

```
val filename = "C:/spark/loudacre_BigData/loudacre_Dataset"
val loudacre_data = sc.textFile(filename)
```

Create a new dataframe having only latitude and longitude

```
val loudacre_data_split = loudacre_data.map(x => x.split(",")).map(s => (s(3),s(4))).toDF
loudacre_data_split.show(10)
```

```
+-----+-----+
|          _1 |          _2 |
+-----+-----+
| 33.6894754264 | -117.543308253 |
| 37.4321088904 | -121.485029632 |
| 39.4378908349 | -120.938978486 |
| 39.3635186767 | -119.400334708 |
| 33.1913581092 | -116.448242643 |
| 33.8343543748 | -117.330000857 |
| 37.3803954321 | -121.840756755 |
| 34.1841062345 |  -117.9435329 |
| 32.2850556785 | -111.819583734 |
| 45.2400522984 | -122.377467861 |
+-----+-----+
only showing top 10 rows
```

- Cleaning

Many of the values are zero in both the columns, customer might have not provided their information about their location. we will remove these rows.

```
loudacre_data_split.filter($"_1"=="0").show(10)
```

```

+---+---+
| _1| _2|
+---+---+
|  0|  0|
|  0|  0|
|  0|  0|
|  0|  0|
|  0|  0|
|  0|  0|
|  0|  0|
|  0|  0|
|  0|  0|
|  0|  0|
+---+---+

```

only showing top 10 rows

```
loudacre_data_split.filter($"_1" === "0").count
```

```
val la_df = loudacre_data_split.filter($"_1" != "0")
```

```
la_df.filter($"_1" === "0").count
```

Convert dataframe back to RDD

```
val la_rdd = la_df.rdd.map(row => List(row.getString(0),row.getString(1)))
```

Create a vector RDD

```
val vectors = la_rdd.map(s => Vectors.dense(s(0).toDouble,s(1).toDouble)).cache()
```

• Train the model

```
val numClusters = 3
```

```
val numIterations = 20
```

```
val kmeansmodel = KMeans.train(vectors,numClusters,numIterations)
```

Display the center points of each cluster

```
kmeansmodel.clusterCenters.foreach(println)
```

```

[34.54436372366799, -118.05513586535038]
[39.92266331983432, -121.38464957351843]
[35.08592000544937, -112.57643826547802]

```

Create a broadcast variable for the Kmeans model

```
kmeansmodel.computeCost(vectors)
```

```
val kmeansModel_BC = sc.broadcast(kmeansmodel)
```

```
val cluster_df = kmeansModel_BC.value.predict(vectors).toDF
```

```
cluster_df.show
```

```
+-----+
|value|
+-----+
|  0 |
|  1 |
|  1 |
|  1 |
|  1 |
|  0 |
|  0 |
|  1 |
|  0 |
|  2 |
|  1 |
|  1 |
|  1 |
|  0 |
|  0 |
|  2 |
|  2 |
|  1 |
|  0 |
|  1 |
|  2 |
+-----+
only showing top 20 rows
```

Create a new column called “id”, which will be used for joining the two dataframes

```
val df1 = cluster_df.withColumn("id",monotonically_increasing_id())
```

```
val df2 = la_df.withColumn("id",monotonically_increasing_id())
```

```
df1.show(5)
```

```
+-----+-----+
|value| id|
+-----+-----+
|  0 |  0 |
|  1 |  1 |
|  1 |  2 |
|  1 |  3 |
|  0 |  4 |
+-----+-----+
only showing top 5 rows
```

```
df2.show(5)
```

```
+-----+-----+-----+
|_1|_2| id|
+-----+-----+-----+
|33.6894754264|-117.543308253| 0|
|37.4321088904|-121.485029632| 1|
|39.4378908349|-120.938978486| 2|
|39.3635186767|-119.400334708| 3|
|33.1913581092|-116.448242643| 4|
+-----+-----+-----+
only showing top 5 rows
```

Join the two dataframes based on “id” column

```
val df3 = df2.join(df1,"id")
```

```
df3.show
```

```
+-----+-----+-----+-----+
|   id |           _1 |           _2 | value |
+-----+-----+-----+-----+
|    26 | 39.4708861702 | -119.659926097 |    1 |
|    29 | 33.8514465953 | -117.787423338 |    0 |
|   474 | 39.5548840745 | -121.020788448 |    1 |
|   964 | 42.8588445306 | -122.490145742 |    1 |
|  1677 | 42.0521070411 | -123.772498927 |    1 |
|  1697 |  34.206121509 | -118.248366666 |    0 |
|  1806 | 36.5039511635 |  -121.0370419 |    1 |
|  1950 | 34.1840832076 |  -118.13214751 |    0 |
|  2040 | 39.4458400243 | -119.337989148 |    1 |
|  2214 | 37.9311384736 | -121.511788497 |    1 |
|  2250 | 34.1648878746 |  -117.67549234 |    0 |
|  2453 | 36.1947181592 | -115.053517304 |    2 |
|  2509 | 38.3183374986 | -120.975396248 |    1 |
|  2529 | 38.1666641293 | -121.845985179 |    1 |
|  2927 | 36.3195030742 | -120.722388463 |    0 |
|  3091 |  37.801310103 |  -122.29490043 |    1 |
|  3506 | 38.3391887988 | -122.403638455 |    1 |
|  3764 | 33.5908087287 | -109.124901699 |    2 |
|  4590 | 37.1075594103 |  -121.72570178 |    1 |
|  4823 | 34.0175895081 |  -111.98383465 |    2 |
+-----+-----+-----+-----+
only showing top 20 rows
```

The number of point for each cluster in the dataset

```
df3.groupBy("value").count().show()
```

```
+-----+-----+
| value | count |
+-----+-----+
|      1 | 177812 |
|      2 |  65686 |
|      0 | 188359 |
+-----+-----+
```

We remove the “id” column and name the other three columns

```
val newNames = Seq("latitude","longitude","cluster")
```

```
val finaldf = df3.drop("id").toDF(newNames:_* )
```

```
finaldf.show
```

| latitude | longitude | cluster |
|---------------|----------------|---------|
| 39.4708861702 | -119.659926097 | 1 |
| 33.8514465953 | -117.787423338 | 0 |
| 39.5548840745 | -121.020788448 | 1 |
| 42.8588445306 | -122.490145742 | 1 |
| 42.0521070411 | -123.772498927 | 1 |
| 34.206121509 | -118.248366666 | 0 |
| 36.5039511635 | -121.0370419 | 1 |
| 34.1840832076 | -118.13214751 | 0 |
| 39.4458400243 | -119.337989148 | 1 |
| 37.9311384736 | -121.511788497 | 1 |
| 34.1648878746 | -117.67549234 | 0 |
| 36.1947181592 | -115.053517304 | 2 |
| 38.3183374986 | -120.975396248 | 1 |
| 38.1666641293 | -121.845985179 | 1 |
| 36.3195030742 | -120.722388463 | 0 |
| 37.801310103 | -122.29490043 | 1 |
| 38.3391887988 | -122.403638455 | 1 |
| 33.5908087287 | -109.124901699 | 2 |
| 37.1075594103 | -121.72570178 | 1 |
| 34.0175895081 | -111.98383465 | 2 |

only showing top 20 rows

- Display the points for each cluster

```
finaldf.filter($"cluster" === "0").show
```

| latitude | longitude | cluster |
|---------------|----------------|---------|
| 33.8514465953 | -117.787423338 | 0 |
| 34.206121509 | -118.248366666 | 0 |
| 34.1840832076 | -118.13214751 | 0 |
| 34.1648878746 | -117.67549234 | 0 |
| 36.3195030742 | -120.722388463 | 0 |
| 34.2298811925 | -117.877102556 | 0 |
| 34.0669923073 | -118.192176947 | 0 |
| 33.0262077564 | -116.812861708 | 0 |
| 34.1521333194 | -116.638723297 | 0 |
| 32.9596188021 | -116.805661125 | 0 |
| 33.9249343611 | -117.881397679 | 0 |
| 33.0239999005 | -116.674524717 | 0 |
| 33.0339156527 | -116.837511787 | 0 |
| 37.038360427 | -119.588745345 | 0 |
| 32.9965920016 | -116.685453769 | 0 |
| 35.6625077365 | -120.285702621 | 0 |
| 34.2486240248 | -118.2386197 | 0 |
| 33.1863998188 | -116.724525124 | 0 |
| 32.9127253094 | -116.307005506 | 0 |
| 33.8233479881 | -117.770690361 | 0 |

only showing top 20 rows

```
finaldf.filter($"cluster" === "1").show
```

| latitude | longitude | cluster |
|---------------|----------------|---------|
| 39.4708861702 | -119.659926097 | 1 |
| 39.5548840745 | -121.020788448 | 1 |
| 42.8588445306 | -122.490145742 | 1 |
| 42.0521070411 | -123.772498927 | 1 |
| 36.5039511635 | -121.0370419 | 1 |
| 39.4458400243 | -119.337989148 | 1 |
| 37.9311384736 | -121.511788497 | 1 |
| 38.3183374986 | -120.975396248 | 1 |
| 38.1666641293 | -121.845985179 | 1 |
| 37.801310103 | -122.29490043 | 1 |
| 38.3391887988 | -122.403638455 | 1 |
| 37.1075594103 | -121.72570178 | 1 |
| 38.4745414188 | -122.134143962 | 1 |
| 37.6543322066 | -121.588940213 | 1 |
| 45.3409177163 | -117.542333377 | 1 |
| 45.1781132987 | -117.661882259 | 1 |
| 42.6990076523 | -122.637552833 | 1 |
| 38.7301629267 | -121.408276173 | 1 |
| 45.4371672468 | -117.700782699 | 1 |
| 37.5033221854 | -121.525720878 | 1 |

only showing top 20 rows

```
finaldf.filter($"cluster" === "2").show
```

| latitude | longitude | cluster |
|---------------|----------------|---------|
| 36.1947181592 | -115.053517304 | 2 |
| 33.5908087287 | -109.124901699 | 2 |
| 34.0175895081 | -111.98383465 | 2 |
| 39.6628949081 | -114.56289137 | 2 |
| 33.6230375795 | -111.542615553 | 2 |
| 35.0789202973 | -111.408292531 | 2 |
| 36.0838304774 | -114.731425185 | 2 |
| 36.2981261276 | -113.613971565 | 2 |
| 32.2734814339 | -111.641629449 | 2 |
| 33.5206745813 | -112.003312523 | 2 |
| 36.661607919 | -115.006223963 | 2 |
| 33.3561047647 | -111.484732167 | 2 |
| 33.3954391971 | -111.861572714 | 2 |
| 32.2866036613 | -111.791915607 | 2 |
| 32.3938041755 | -111.506443379 | 2 |
| 33.6603116416 | -111.513103129 | 2 |
| 35.0455801515 | -111.316399941 | 2 |
| 33.7285608333 | -111.825495401 | 2 |
| 32.5350707596 | -110.614014808 | 2 |
| 33.6029085259 | -111.304704398 | 2 |

only showing top 20 rows